

Contents

Proposal: Multi-Engine Classroom (conceptual)

Why this exists
Vision cite (this is already documented)
Why multi-engine (Koen's three reasons)
Initial 3-engine setup (concrete sketch)
Initial scenario list (for discussion)
Markov usage graph (conceptual)
Sources consulted
Out of scope for now
Ask of Koen

Proposal: Multi-Engine Classroom (conceptual)

Status: Conceptual draft for discussion — supersedes [multi-engine-classroom-scenarios.md](#) (2026-09-28).

Revision: 2026-09-30b — pointer to separate Operator Console Markov ([multi-engine-operator.md](#)). Usage Markov (Kolibri / Nextcloud / Wikipedia) unchanged.

Author: Steve (Lead Bot), 2026-09-29 (rev. 2026-09-30b)

Audience: Koen / IDEA leads

Companion capacity issue: [idea#159](#) — measure safe concurrent Kolibri video streams per instance

Why this exists

Koen asked for a clean conceptual framing of **why a school runs more than one Engine**, how that looks with a small named fleet, what classroom scenarios matter, and how a Markov usage graph can drive realistic tests later.

This document stays **conceptual**. Implementation, discovery protocol, auth, Playwright harnesses, fleet claim scripts, and Engine code changes are explicitly out of scope here (see final section).

Vision cite (this is already documented)

Prior drafting treated multi-Engine load distribution as an undocumented assumption. That was wrong.

It is already in the IDEA vision / solution description:

All Appdockers on the same LAN will automatically find one another and when a new Appdocker device is added, the catalog of available apps is automatically extended with the apps docked onto the new device.
Performance is optimized by adding Appdockers and redistributing the apps over the Appdockers.

Source: [agent-engine-dev/proposals/solution-description.md](#)

(local clones also under e.g. `/workspace/research/agent-engine-dev/proposals/solution-description.md`)

Referenced from `idea/docs/grok-bot-setup.md` as the Engine vision and intent doc.

The multi-Engine classroom story below **extends** that vision into classroom roles and usage scenarios; it does not invent the load-spread principle.

Why multi-engine (Koen's three reasons)

1. Distribute apps / load across Engines

The physical act is simple: **add an Engine** and **plug a demanding app into it**.

- Engines **autofind and connect** on the school LAN (vision: Appdockers find one another).
- The **integrated app list** looks the same whether the school has 1 Engine or N Engines.
- Users **never need to know** which Engine serves which app, or which Console they opened — any Console on any Engine can see the full picture.

This is the same physical metaphor as the rest of IDEA: capacity is improved by adding hardware and redistributing disks, not by asking teachers to become sysadmins.

2. Scale one app across Engines

Some apps hit a **per-instance / per-Pi limit**. Example: concurrent Kolibri video streams.

When one Kolibri instance cannot safely serve the whole class:

- Start a **second Kolibri instance** on another Engine.

- **Assign students** across instances (class A → instance 1, class B or half the cohort → instance 2).
- Same pattern later for Nextcloud video / heavy media.

Capacity baseline for Kolibri ([idea#159](#), idea01 Pi 5 / 4 GB). Multi-quality **planning** streams per instance (LAN measured SAFE ≥128 for all four; planning leaves Wi-Fi / browser / Pi-4 headroom):

Quality	Resolution	Planning streams / instance
low	360p (640×360)	64
medium	480p (854×480)	32
prior / typical lesson	720p (1280×720)	32
high	1080p (1920×1080)	12

Default classroom planning figure remains **32 @ ~720p**; use 64 / 12 when content is known low / high. Load-test tooling: [agent-app-dev#9](#).

3. Per-class Engine with own instances — Console manages everything from anywhere

Each class (or teaching block) can have **its own Engine** hosting **class-only** app instances — e.g. Class 5A Nextcloud, Class 5A Kolibri content — without isolating operators.

Any Console on any Engine can manage **all apps across all Engines** simply. Teachers and operators are not trapped on “their” Pi’s Console; the mesh presents one school.

Initial 3-engine setup (concrete sketch)

Design target first: a **named three-Engine school**, then generalise to N.

Name	Role (story)	Typical disks / instances
idea-A	Shared / school golden hub	Shared apps everyone uses lightly: offline Wikipedia (Kiwix), school-wide Nextcloud “School Files”, maybe a small shared Kolibri channel library. Stable reference Engine.
idea-B	Class Engine — Grade 5A	Class-only Kolibri (Grade 5A lessons/quizzes), class Nextcloud (Group “Grade 5A”: view-only shares, File Drop, collaborative docs).
idea-C	Class Engine — Form 3	Class-only Kolibri (Form 3), class Nextcloud (Group “Form 3”). When Grade 5A video load overflows idea-B, a second Kolibri instance for 5A can also land here (reason 2).

Classroom story: Monde Primary (or any Marco site) starts with one Engine in the room. When video buffering and “limit to one group at a time” become the norm (field troubleshooting already warns about many devices), the school adds **idea-B** and **idea-C** as classroom Engines. Teachers still open any Console; students still join `appnet` and pick apps from one list. Operators dock demanding class disks onto B/C and leave shared services on A.

Generalising to N

The three-Engine model is the **design target first**. N Engines later means:

- More class Engines (idea-D...), and/or
- More replicas of a hot app (extra Kolibri / Nextcloud instances) assigned by cohort.

Discovery, Console overview, and “users don’t pick an Engine” stay the same rules; only inventory and assignment grow.

Initial scenario list (for discussion)

Discussable outcomes — **not** test scripts. Named UI scripts that expand Markov nodes live under [Detailed UI scenarios](#).

Reason 1 — distribute apps / load

- **S1 Add Engine, plug demanding app:** School adds idea-C; docks a heavy Kolibri (or Nextcloud media) disk onto it; app appears in every Console’s integrated list without renaming Consoles or teaching users a new hostname.

- **S2 Redistribute load:** Move or copy a demanding instance from idea-A onto idea-B so A stays responsive for shared services; app list and Open links remain coherent school-wide.
- **S3 Invisible Engine identity:** Student or teacher opens Console on idea-B or idea-C and sees the same app catalog/status as on idea-A.

Reason 2 — scale one app

- **S4 Kolibri stream ceiling:** Class exceeds safe concurrent streams on one instance → second Kolibri on another Engine; students assigned so neither instance exceeds the [idea#159](#) planning N for that video quality (**64 @ 360p / 32 @ 480p–720p / 12 @ 1080p**).
- **S5 Nextcloud media twin:** Same idea for Nextcloud video / large file concurrent access — second instance + cohort assignment when one Pi is the bottleneck.

Reason 3 — per-class Engine + Console-anywhere

- **S6 Class-only content:** Grade 5A Kolibri/Nextcloud live only on idea-B; Form 3 only on idea-C; shared Wikipedia (Kiwix) stays on idea-A.
- **S7 Manage from the other room:** Teacher or operator on idea-A's Console starts/stops, checks status, or prepares disks for apps that run on idea-B and idea-C.
- **S8 Cross-class share without moving rooms:** Teacher on idea-B Console shares a Nextcloud folder with Form 3's group whose home instance is on idea-C (management is school-wide; data placement follows class Engines).

Usage / field-aligned (Marco)

- **S9 Kolibri teaching cycle:** Teacher builds class → enrolls learners → assigns lesson with videos → students watch/complete → quiz → teacher reads Reports (from Marco Kolibri presentations). Scripts: `K-manage-*` , `K2` , `K4` .
- **S10 Nextcloud classroom workflow:** Class groups; view-only material share; File Drop homework; collaborative doc; Talk for support (from Marco Nextcloud presentations). Scripts: `NC-share` , `NC-drop` , `NC-collab` , `NC1` / `NC2` . Field note: many devices → performance drop → prefer one group at a time unless multi-Engine capacity is in play.
- **S11 Offline Wikipedia:** Learner (or teacher) opens Kiwix from Console (idea-A), searches/browses articles, dwells, leaves. Script: `W1` .

Markov usage graph (conceptual)

Idea

- **Nodes** = usage states (what a person is doing in an app or Console), not disk-dock hardware states alone.
- **Transitions** = actions with probabilities (initial guesses until field frequencies exist). Compact **scenario IDs** on edges (e.g. `[L-entry]` , `[K-manage-*`] , `[W1]`) mark where a transition is the **start** of a named UI script; the harness expands the destination node into ordered clicks.
- A **real test scenario** = a **random walk** on this graph for some duration or step count.
- Test instances of each app ship with **preloaded content**; legal actions are derived from that content (e.g. which lessons/videos exist, which folders/groups exist).
- Incorporate Marco's file-sharing (students/teachers) and Kolibri classroom management, plus content-access actions for Kolibri, Nextcloud, and offline Wikipedia (Kiwix on idea-A).

Operator Console Markov (separate): managing/altering the multi-Engine setup (login, NetworkTree, dock/eject, install, start/stop, copy/move, Files/Backup/erase, operators, settings) lives in sibling [multi-engine-operator.md](#) with graph files `multi-engine-operator-markov.{dot,png,svg}` . This classroom usage graph is unchanged aside from this pointer.


Example nodes (small set)

Node	Who	Colour	Meaning	Scenario IDs that expand it
<code>console_overview</code>	Teacher / operator	blue	School-wide Engine/app status; opens apps as coach	entry for <code>K-manage-*</code> , <code>NC1</code> , <code>W1</code> (teacher)
<code>console_learner</code>	Student	mid-blue	Same unified app list; student opens an app to work	<code>L-entry</code>

Node	Who	Colour	Meaning	Scenario IDs that expand it
<code>kolibri_home</code>	Student or teacher	green	Kolibri signed in, Learn	K2 (learner start of lesson)
<code>kolibri_watching</code>	Student	green	Playing an assigned lesson video	K2 , K3 (capacity dwell)
<code>kolibri_exercise</code>	Student	green	Doing an exercise in a lesson	K2
<code>kolibri_manage</code>	Teacher	green (dark)	Classes, lessons, quizzes, Reports	K-manage-class , K-manage-enroll , K-manage-lesson , K-manage-quiz , K-manage-reports (+ composite K1 / K4)
<code>nc_browse</code>	Student or teacher	yellow	Nextcloud Files browse	hub for NC-share / NC-drop / NC-collab / NC1
<code>nc_share</code>	Teacher	yellow (dark)	View-only share to a class group	NC-share
<code>nc_drop</code>	Student	yellow (dark)	Uploading to a File Drop	NC-drop
<code>nc_collab</code>	Student	yellow (dark)	Editing a shared collaborative doc	NC-collab
<code>wiki_browse</code>	Student or teacher	purple	Kiwix / offline Wikipedia on idea-A	W1

Example transitions (probabilities are placeholders)

Graphviz source: [multi-engine-markov.dot](#) . Rendered:



► Same graph as Mermaid (editable)

Story walkthrough (nodes → real clicks)

How this maps to tests. Each Markov **node** is a coarse place someone can be (learner Console, teacher Console, Kolibri coaching, watching a video, Wikipedia, ...). A **transition** is leaving that place for another. Inside a node, a real test does **not** invent a new node for every mouse click — that would explode the graph. Instead the harness expands the node into an **ordered UI script** (login → sidebar → form → Save) taken from Marco's classroom presentations. Create an extra node only when the action changes **who** is acting, **which app surface** they are on, or **resource pressure** (e.g. video stream vs browse). Probabilities below are still placeholders. Scenario IDs in [brackets] on edges/nodes name the scripts under [Detailed UI scenarios](#).

UI labels below match Marco's field decks and quick-reference cards (kolibri-classroom-setup , kolibri-lessons-and-quizzes , nextcloud-user-registration , nextcloud-classroom-use , quick-reference-kolibri.md , quick-reference-nextcloud.md) — confirmed with Marco 2026-09-29. Console steps assume the school-wide app overview any Engine's Console shows (reason 1 / 3). Exact selectors belong in the Playwright harness later — this doc names the human-visible controls.

console_overview — teacher / operator opens the Console

The teacher or operator opens a browser on the school LAN (or already has Console open) and lands on the **Engine / apps overview**: Engines present, apps docked, status. They are inspecting what is available on the network, not choosing a "home" Pi. **Operator-specific Console actions** (dock/eject disks, start/stop instances, install/move/backup/erase, operators/settings) are **not** on this usage graph — see the separate Operator Console Markov: [multi-engine-operator.md](#) (01 – 018 , multi-engine-operator-markov.*).

Concrete UI (when a walk starts here):

1. Open Chromium (or any browser) on the school LAN → Console for any Engine (field habit today: hostname like engine-1.local ; multi-Engine story: idea-A / idea-B / idea-C Consoles must be equivalent).
2. Scan the integrated app list — apps show **Running** (etc.) — and Engine rows.
3. Optionally refresh / wait while students work.

Scenario IDs started from here: K-manage-* (open Kolibri as teacher), NC1 / NC-share (open Nextcloud), W1 (open Wikipedia as teacher), plus dwell.

From here the walk can:

1. **Open Kolibri as teacher** (~0.20) [K-manage-*] — click the Kolibri app (Running) → **Sign in** on Kolibri → go to [kolibri_manage](#) .
2. **Open Nextcloud** (~0.20) [NC1 / NC-share] — click the Nextcloud app → Log in → go to [nc_browse](#) .
3. **Open Wikipedia / Kiwix** (~0.10) [W1] — click the Kiwix / Wikipedia app (on idea-A) → go to [wiki_browse](#) .
4. **Stay on overview** (~0.50) — refresh / dwell → remain on **console_overview** .

Test note: opening an app is one transition; the login form that follows belongs to the destination node's script, not a separate Console node.

console_learner — student opens the Console

The **learner** entry that was missing from the earlier draft. A student opens Console, sees the **same unified app list** as teachers (reason 1 / S3), and opens an app to start work. They are not managing Engines — just picking Kolibri, Nextcloud, or Wikipedia.

Concrete UI (walk starts here — scenario L-entry):

1. Open browser on school LAN → any Engine's Console (idea-A / B / C equivalent).
2. See integrated app list (Kolibri class instance, Nextcloud class instance, Kiwix on idea-A, ...).
3. Click one app → destination node's login / home script runs.

Scenario IDs: L-entry (this node + open-app transitions); destination scripts K2 / NC-drop / NC-collab / W1 continue after landing.

From here the walk can:

1. **Open Kolibri** (~0.40) [L-entry] → [kolibri_home](#) .
2. **Open Nextcloud** (~0.25) [L-entry] → [nc_browse](#) .
3. **Open Wikipedia / Kiwix** (~0.20) [L-entry / W1] → [wiki_browse](#) .
4. **Dwell** (~0.15) → remain on **console_learner** .

kolibri_manage — teacher in Kolibri coaching

The teacher is on Kolibri's **facility / coaching** side (Classes, Lessons, Quizzes, Reports) — not the learner Learn tab. This node **generalises** several coaching tasks; a single visit expands into **one** of the K-manage-* scripts below (create class, enroll, build lesson, quiz, reports). The graph does **not** add nodes for "click Learners tab" vs "click Save" — those are steps inside the script.

Entry UI (always, if not already signed in):

1. Kolibri login page → type **teacher username** and **password** → Sign in.
2. Land on teacher home / facility UI (left sidebar visible).

Typical actions inside this node (pick one path per visit; stay in-node until done) — each row = one named scenario:

Intent	Clicks (Marco)	Scenario ID
Create class	Left sidebar Classes → + New class → type name (e.g. Grade 5A) → Save	K-manage-class
Enroll learners	Open class → Learners tab → Enroll learners → tick students → Confirm	K-manage-enroll
Build / enrich lesson	Class → Lessons → + New lesson (or open existing) → name → Add resources → pick video + exercise from Library / Channels → Save → set Recipients to the class → toggle Visible	K-manage-lesson
Create quiz	Class → Quizzes → + New quiz → Add questions from exercise channels → set count → Finish → toggle Active	K-manage-quiz
Read progress	Left sidebar Reports → Classes → class → Lessons or Quizzes → open item → scan learner table	K-manage-reports

Scenario IDs that expand this node: K-manage-class, K-manage-enroll, K-manage-lesson, K-manage-quiz, K-manage-reports.
Composite walks: K1 (class+enroll+lesson), K4 (quiz+reports with learner take).

From here the walk can leave to:

1. **Preview as learner** (~0.20) — open learner view / Learn (or sign in as a test learner) → [kolibri_home](#).
2. **Keep coaching** (~0.55) [K-manage-*] — another coaching intent from the table → stay on [kolibri_manage](#).
3. **Back to Console** (~0.25) — leave Kolibri tab / return to Console URL → [console_overview](#).

[kolibri_home](#) — learner on Kolibri Learn

A student (or teacher previewing) is on the **learner** side after login.

Entry UI:

1. Kolibri login → **learner username** / **password** → Sign in. (Reached via [L-entry](#) from [console_learner](#), or teacher preview from [kolibri_manage](#).)
2. Top menu **Learn**.
3. See enrolled classes → open class → see assigned **Lessons** (and **Quizzes** if active).

Scenario IDs: K2 (open video / exercise); learner half of K4 (open quiz).

From here:

1. **Open a lesson video** (~0.45) [K2] — open lesson → click the **video** resource → player starts → [kolibri_watching](#).
2. **Open an exercise** (~0.25) [K2] — open lesson → click an **exercise** resource → [kolibri_exercise](#).
3. **Browse only** (~0.20) — move between classes / lessons without starting a resource → stay on [kolibri_home](#).
4. **Leave** (~0.10) → [console_learner](#) (or close app).

[kolibri_watching](#) — student in the video player

Capacity-sensitive node (reason 2 / idea#159). Many concurrent walkers here is what forces a second Kolibri instance. Planning N depends on video quality: **64 @ 360p / 32 @ 480p-720p / 12 @ 1080p** (idea01 Pi 5 / 4 GB).

UI while here: video playing in-browser (no download); student may pause/seek; Kolibri records progress in the background.

Scenario IDs: K2 (playthrough), K3 (N walkers dwell for stream ceiling).

From here:

1. **Keep watching** (~0.60) [K3] → stay on [kolibri_watching](#).
2. **Next resource** (~0.25) — leave player, open next lesson item (often exercise) → [kolibri_exercise](#).

3. **Exit lesson** (~0.15) — back to class / Learn → [kolibri_home](#) .

kolibri_exercise — student in an exercise

UI: exercise questions with immediate feedback; student answers and submits; Kolibri records started / completed / score.

Scenario IDs: `k2` (exercise half of lesson).

From here:

1. **Finish / exit** (~0.70) → [kolibri_home](#) .
2. **Next video** (~0.30) → [kolibri_watching](#) .

(Quiz-taking can reuse this node or stay inside `kolibri_home` → open Quizzes → Start → answer → submit; if quizzes become a distinct load profile later, split a `kolibri_quiz` node then. Mapped in scenario `k4` / `k-manage-quiz` .)

nc_browse — anyone in Nextcloud Files

Hub for Marco's classroom file workflow. **Files** app is the central UI (manual p.12 in Marco's deck).

Entry UI:

1. Nextcloud login → username / password (teacher via `console_overview` , student via `L-entry` / `console_learner`).
2. Open **Files** (default home).
3. Browse folders (class materials, Drop Zone, school shares).

Scenario IDs started from here: `NC-share` , `NC-drop` , `NC-collab` (and composite `NC1`).

From here:

1. **Share to class** (teacher, ~0.18) [`NC-share`] → [nc_share](#) .
2. **File Drop** (student on drop link, or teacher preparing drop, ~0.22) [`NC-drop`] → [nc_drop](#) .
3. **Open collab doc** (~0.22) [`NC-collab`] → [nc_collab](#) .
4. **Keep browsing** (~0.15) → stay on `nc_browse` .
5. **Leave as learner** (~0.13) → [console_learner](#) .
6. **Leave as teacher** (~0.10) → [console_overview](#) .

nc_share — teacher sharing a folder

Matches table / scenario `NC-share` (view-only materials row of the Nextcloud classroom workflow).

UI script (view-only materials):

1. In Files, hover the folder (e.g. `Videos to watch`).
2. Click **Share** (person + icon).
3. Under **Internal shares**, type class group name (e.g. `Class2A` / `Grade 5A`) → select group.
4. Set permission **View only**.

Group must already exist (Accounts → Groups → + — see scenario `NC1` preload; that admin path can stay a preload step, not a graph node).

From here:

1. **Done sharing** (~0.80) → [nc_browse](#) .
2. **Leave** (~0.20) → [console_overview](#) .

nc_drop — File Drop upload

Matches scenario `NC-drop` .

Teacher prep (often preload, or expand inside a long `nc_browse` visit): Share icon on `Drop Zone` → + beside Create public link → **File request** → copy link (clipboard) → optionally paste into Talk.

Student UI (this node):

1. Open the File Drop / file-request link.

2. See empty “Click or drop to upload” UI (cannot see others' files).
3. Choose file(s) / drop → upload completes.

From here:

1. **After upload** (~0.85) → [nc_browse](#) (or close link).
2. **Leave** (~0.15) → [console_learner](#) .

nc_collab — live document editing

Matches scenario **NC-collab** .

UI:

1. (Teacher create, often preload:) Files → **+ New** → **New Document** → name → pick template → Share with group → **Allow editing**.
2. Student: open shared file from Files → editor loads; peer **avatars** top-right; cursors move live.

From here:

1. **Close doc** (~0.70) → [nc_browse](#) .
2. **Keep editing** (~0.30) → stay on **nc_collab** .

wiki_browse — Kiwix / offline Wikipedia

Shared service on **idea-A** (school golden hub). Learners and teachers reach it from Console (**L-entry** / **W1** from **console_learner** , or **W1** from **console_overview**). Light load relative to Kolibri video — still a first-class classroom usage state.

Entry UI:

1. From Console, click **Kiwix** / Wikipedia app (Running on idea-A).
2. Kiwix library / Wikipedia ZIM opens in the browser.
3. Use search box or browse categories / random article.

Scenario IDs: **W1** .

From here:

1. **Search / browse** (~0.65) [**W1**] — type a query, open articles, follow links, dwell reading → stay on **wiki_browse** .
2. **Leave as learner** (~0.20) → [console_learner](#) .
3. **Leave as teacher** (~0.15) → [console_overview](#) .

Detailed UI scenarios (expand nodes into clicks)

These are **test scripts**, not extra Markov nodes. Each script names the Markov node(s) it expands and, where relevant, the matching table row. A walker that lands on **kolibri_manage** for “build lesson” runs **K-manage-lesson** ; N concurrent **kolibri_watching** walkers are **K3** (capacity). Preload content so every click target exists.

Mapping summary

Scenario ID	Markov node(s)	Matches
L-entry	console_learner → kolibri_home / nc_browse / wiki_browse	Learner opens Console, picks an app
K-manage-class	kolibri_manage	Manage table: Create class
K-manage-enroll	kolibri_manage	Manage table: Enroll learners
K-manage-lesson	kolibri_manage	Manage table: Build / enrich lesson
K-manage-quiz	kolibri_manage	Manage table: Create quiz
K-manage-reports	kolibri_manage	Manage table: Read progress

Scenario ID	Markov node(s)	Matches
K1	console_overview → kolibri_manage (chains class+enroll+lesson)	Composite of three manage rows
K2	console_learner → kolibri_home → kolibri_watching / kolibri_exercise	Student completes assigned lesson
K3	N× kolibri_watching dwell	Stream ceiling / S4 / idea#159
K4	kolibri_manage (K-manage-quiz + K-manage-reports) + learner quiz take	Quiz + Reports
NC-share	nc_browse → nc_share	View-only share to class group
NC-drop	nc_browse → nc_drop	File Drop upload
NC-collab	nc_browse → nc_collab	Collaborative doc edit
NC1	console_overview → nc_browse + share/drop/collab	Composite S10 teacher → students
NC2	console_overview (idea-A) → class NC on idea-B	Manage from other room / S7
W1	console_* → wiki_browse	Open Kiwix, search/browse, leave
C1	console_overview dwell / open-app	Add Engine, app appears everywhere / S1

Scenario L-entry — Student opens Console and starts work

Role: student · **Engines:** any Console; destination app may be on idea-A (Kiwix), idea-B (class Kolibri/Nextcloud), etc.

Markov: * → console_learner → (kolibri_home | nc_browse | wiki_browse).

1. Browser → Console on school LAN (idea-A / B / C — same catalog).
2. Scan unified app list (Running apps).
3. Click **Kolibri** → learner login → land on kolibri_home (continue with K2), or
4. Click **Nextcloud** → student login → land on nc_browse (continue with NC-drop / NC-collab), or
5. Click **Kiwix / Wikipedia** → land on wiki_browse (continue with W1).

Scenario K-manage-class — Create class

Role: teacher · **Markov:** console_overview → kolibri_manage (self-loop / visit).

Table row: Create class.

1. Console → Open Kolibri (class instance on idea-B) → teacher login → kolibri_manage .
2. Left sidebar **Classes** → **+ New class** → type name (e.g. Grade 5A) → **Save**.
3. Stay on kolibri_manage or back to console_overview .

Scenario K-manage-enroll — Enroll learners

Role: teacher · **Markov:** kolibri_manage .

Table row: Enroll learners. *Requires class (preload or K-manage-class).*

1. Open class → **Learners** tab → **Enroll learners** → tick students → **Confirm**.
2. Stay on kolibri_manage or leave to Console.

Scenario K-manage-lesson — Build lesson + assign

Role: teacher · **Markov:** kolibri_manage .

Table row: Build / enrich lesson.

1. Class → **Lessons** → **+ New lesson** (or open existing) → name (e.g. Week 1: Fractions Introduction).
2. **Add resources** → Library/Channels → select **1 video** + **1 exercise** → **Save**.
3. Confirm **Recipients** = class → toggle lesson **Visible**.
4. Stay on kolibri_manage , preview as learner → kolibri_home , or back to Console.

Scenario `K-manage-quiz` — New quiz

Role: teacher · **Markov:** `kolibri_manage` .

Table row: Create quiz.

1. Class → **Quizzes** → + **New quiz** → **Add questions** from exercise channel → set count (e.g. 5–10) → **Finish** → toggle **Active**.
2. Stay on `kolibri_manage` (often followed by `K-manage-reports` after learners take it — see `K4`).

Scenario `K-manage-reports` — Read Reports

Role: teacher · **Markov:** `kolibri_manage` .

Table row: Read progress.

1. Left sidebar **Reports** → **Classes** → class → **Lessons** or **Quizzes** → open item → scan learner table (scores / completion).
2. Stay on `kolibri_manage` or back to `console_overview` .

Scenario `K1` — Teacher prepares Grade 5A lesson (composite)

Role: teacher · **Engines:** Console on idea-A or idea-B; Kolibri on idea-B.

Markov: `console_overview` → `kolibri_manage` (chains `K-manage-class` → `K-manage-enroll` → `K-manage-lesson` , optional `K-manage-reports`).

1. Console: open Kolibri for Grade 5A → transition `console_overview` → `kolibri_manage` .
2. Run `K-manage-class` (*skip if class preloaded*).
3. Run `K-manage-enroll` .
4. Run `K-manage-lesson` (video + exercise, Visible).
5. Optional: `K-manage-reports` (*empty progress yet*).
6. Either stay (`kolibri_manage` self-loop) or **preview as learner** → `kolibri_home` , or back to Console.

Scenario `K2` — Student completes assigned lesson

Role: student · **Engine:** same Kolibri as teacher prep.

Markov: `console_learner` [L-entry] → `kolibri_home` → `kolibri_watching` / `kolibri_exercise` .

1. L-entry : Console → Open Kolibri → login learner credentials → `kolibri_home` .
2. **Learn** → open **Grade 5A** → open assigned lesson.
3. Click video resource → `kolibri_watching` (play through or paced load-test equivalent).
4. Click exercise resource → `kolibri_exercise` → answer items → submit.
5. Return to Learn / class → `kolibri_home` , or leave to `console_learner` .

Scenario `K3` — Class hits stream ceiling (reason 2 / S4)

Setup: lesson with a video pre-visible; idea#159 planning N by quality — **64 @ 360p / 32 @ 480p–720p / 12 @ 1080p** (idea01 Pi 5 / 4 GB).

Markov: N walkers dwell in `kolibri_watching` (still one node — concurrency is walker count).

1. N student walkers each run `K2` steps 1–3 and **dwell** in `kolibri_watching` .
2. If N > planning N for that quality on idea-B's Kolibri: start second Kolibri on idea-C; assign half the walkers to each instance (same lesson content shape).
3. Console on idea-A still lists both apps; teachers do not pick Engine by hostname (S3).

Scenario `K4` — Quiz + Reports

Markov: `kolibri_manage` (`K-manage-quiz` + later `K-manage-reports`) + short learner visit on `kolibri_home` / exercise-like load.

1. Teacher: run `K-manage-quiz` (Active).
2. Students (`kolibri_home` via L-entry): class → **Quizzes** → **Start** → answer → submit (*treat as exercise-like load or keep under home until a quiz node is justified*).
3. Teacher: run `K-manage-reports` on that quiz → read scores / drill into a learner.

Scenario `NC-share` — View-only share to class group

Role: teacher · **Markov:** nc_browse → nc_share → nc_browse .

Matches: Nextcloud classroom “share materials view-only” row.

1. In Files, hover folder (e.g. Videos to watch) → **Share**.
2. Internal share to class group (e.g. Grade 5A) → permission **View only**.
3. Done → back to nc_browse (or leave to console_overview).

Scenario NC-drop — File Drop homework upload

Role: student (teacher prep optional) · **Markov:** nc_browse → nc_drop → nc_browse / console_learner .

Matches: File Drop / file-request row.

1. (Teacher prep / preload:) Drop Zone → Share → public link + → **File request** → copy link.
2. Student: open File Drop link → upload file(s) → nc_drop .
3. After upload → nc_browse or leave to console_learner .

Scenario NC-collab — Collaborative document

Role: student (teacher create often preload) · **Markov:** nc_browse → nc_collab .

Matches: collaborative doc row.

1. (Teacher create / preload:) + **New** → **New Document** → share group with **Allow editing**.
2. Student: open shared file from Files → edit live (nc_collab) → close → nc_browse .

Scenario NC1 — Teacher share + File Drop + collab (S10 composite)

Role: teacher then students · **Engine:** class Nextcloud on idea-B.

Markov: console_overview → nc_browse + NC-share / NC-drop / NC-collab .

1. Console → Open Nextcloud → nc_browse (teacher login).
2. (Preload or once) Profile → **Accounts** → Groups + → Grade 5A ; + **New account** students into that group.
3. Run NC-share on Videos to watch .
4. Prepare File Drop on Drop Zone (Talk optional: top menu Talk → + conversation → share link).
5. Create collab doc + share with editing.
6. Students via L-entry : open shared folder (nc_browse); run NC-drop ; run NC-collab .

Scenario NC2 — Manage class Nextcloud from another room (S7)

Markov: console_overview on idea-A while class NC runs on idea-B.

1. Teacher opens **Console on idea-A** (console_overview) while Grade 5A Nextcloud runs on idea-B.
2. Open that Nextcloud from the integrated list → same NC1 / NC-share steps.
3. Assert: management works without walking to idea-B's physical Console (reason 3).

Scenario W1 — Browse offline Wikipedia (Kiwix on idea-A)

Role: learner or teacher · **Engine:** Kiwix on idea-A.

Markov: console_learner or console_overview → wiki_browse (self-loop) → leave to matching Console.

1. Console → Open **Kiwix / Wikipedia** → wiki_browse .
2. Search box: type a topic (e.g. class-related keyword) → open an article.
3. Follow internal links / browse categories; dwell reading (self-loop).
4. Leave → console_learner (student) or console_overview (teacher).

Scenario C1 — Add Engine, app appears everywhere (S1)

1. Operator on idea-A Console: note app list. (Operator dock/start actions are not modelled in this usage graph — see [multi-engine-operator.md](#) scenarios 04 / 09 / 06 .)
2. idea-C comes online with a docked Kolibri/Nextcloud (hardware/Ops outside Markov).
3. Refresh Console on idea-A **and** idea-B **and** a learner Console view: new app row present → Open works → enters kolibri_manage / kolibri_home / nc_browse as usual.

4. No new Markov nodes — still `console_overview` / `console_learner` dwell / open-app transitions; assertion is catalog equality across Consoles.

How a walk becomes a test

- 1. Preload Kolibri with a class, enrolled learners, and a lesson that includes at least one video and one exercise (Marco Week 1–2 content shape).
- 2. Preload Nextcloud with class groups, a view-only materials folder, a File Drop folder, and one collaborative document (Marco classroom-use shape).
- 3. Preload / confirm Kiwix ZIM available on idea-A.
- 4. Sample a walk: e.g. `console_learner` → `kolibri_home` → ... → `kolibri_watching` (N concurrent walkers = N “students” hitting video — bounded by idea#159 quality table), or `console_overview` → `kolibri_manage` for coaching scripts.
- 5. Multi-Engine walks assign walkers to instances on idea-B vs idea-C when testing reason 2; Console actions may start on any Engine's Console (reason 3).

Full YAML / runner design stays with the existing Markov duration-tests proposal (`duration-tests.md` , idea#85) when implementation starts — not here.

Sources consulted

Source	Use
<code>agent-engine-dev/proposals/solution-description.md</code>	Vision: autofind Appdockers; add Appdockers + redistribute apps for performance
<code>multi-engine-classroom-scenarios.md</code> (+ .pdf)	Prior draft — superseded (implementation-heavy; incorrectly framed multi-Engine as undocumented)
Marco / programme-manager: <code>presentations/kolibri-classroom-setup</code> , <code>kolibri-lessons-and-quizzes</code> , <code>nextcloud-user-registration</code> , <code>nextcloud-classroom-use</code> ; <code>reference-cards/quick-reference-kolibri.md</code> , <code>quick-reference-nextcloud.md</code>	Exact UI labels for Markov story + scenarios (Classes, Learn, Share, File request, ...)
Marco field: <code>field/troubleshooting-guide.md</code> (via prior draft + local copy)	Many devices → performance drop; limit one group at a time
<code>duration-tests.md</code>	Existing Markov duration-test design (implementation later)
idea#159 (+ multi-quality comment)	Kolibri concurrent stream capacity — planning 64/32/32/12 by quality

Out of scope for now

Flagged for later proposals / issues — **not** designed here:

- Implementation of multi-Engine discovery / peering changes
- Auth, operator roles, and student-assignment UX between Kolibri instances
- Playwright / Console e2e harness, fleet claim scripts, fixture disks
- Exact Engine APIs for “assign students across instances”
- Hardware power-cut, USB disk carry automation, golden-Pi rules
- PDF/print packaging of field guides
- ~~Operator Console Markov graph~~ — **done as sibling**: `multi-engine-operator.md` + `multi-engine-operator-markov.*` (kept separate from this usage graph)

When Koen agrees the conceptual shape, Steve can split follow-ups (Engine / Console / App Dev / Ops) without baking premature design into this doc.

Ask of Koen

1. Confirm the **three reasons** match intent.
2. Confirm **idea-A / idea-B / idea-C** as the named 3-Engine story (or rename).
3. Confirm **learner Console** (`console_learner`) + **Wikipedia** (`wiki_browse`) + **K-manage-*** ↔ **manage table** wiring.
4. Prioritise which of S1–S11 to deepen next; review the separate Operator Console Markov ([multi-engine-operator.md](#)).
5. Treat idea#159 quality table (64/32/32/12) as the capacity gate for reason-2 scenarios.